

**Assignment -2**  
Python Programming

|                     |                   |
|---------------------|-------------------|
| Assignment Date     | 23 september 2022 |
| Student Name        | MATHIVATHANI.B.G  |
| Student Roll Number | 311419205021      |
| Maximum Marks       | 2 Marks           |

# Data Visualization and Pre-processing

**Question-1:** 1.

Load the dataset

**Solution:**

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
sns.set_style('darkgrid')
sns.set(font_scale=1.3)
```

In [2]:

```
df=pd.read_excel("/content/Churn_Modelling.xlsx")
```

```
In [1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
sns.set_style('darkgrid')
sns.set(font_scale=1.3)
```

```
In [2]: df=pd.read_excel("/content/Churn_Modelling.xlsx")
```

**Question-2:**

2. Perform Below Visualizations.


- Univariate Analysis
- Bi - Variate Analysis ●
- Multi - Variate Analysis

**Solution:**


```
.
#Perform Univariate Analysis
plt.figure(figsize=(8,8))
sns.kdeplot(x=df['Balance'])
```

```
In [7]: #Perform Univariate Analysis
plt.figure(figsize=(8,8))
sns.kdeplot(x=df['Balance'])
```

```
Out[7]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc3f3579c50>
```



```
#Perform Bivariate Analysis plt.scatter(df.Age,df.Balance)
```



```
#Perform Bivariate Analysis df.corr()
```

```
Out[9]:
```

|                 | CreditScore | Age       | Tenure    | Balance   | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSalary | Exited    |
|-----------------|-------------|-----------|-----------|-----------|---------------|-----------|----------------|-----------------|-----------|
| CreditScore     | 1.000000    | -0.003965 | 0.000842  | 0.006268  | 0.012238      | -0.005458 | 0.025651       | -0.001384       | -0.027094 |
| Age             | -0.003965   | 1.000000  | -0.009997 | 0.028308  | -0.030680     | -0.011721 | 0.085472       | -0.007201       | 0.285323  |
| Tenure          | 0.000842    | -0.009997 | 1.000000  | -0.012254 | 0.013444      | 0.022583  | -0.028362      | 0.007784        | -0.014001 |
| Balance         | 0.006268    | 0.028308  | -0.012254 | 1.000000  | -0.304180     | -0.014858 | -0.010084      | 0.012797        | 0.118533  |
| NumOfProducts   | 0.012238    | -0.030680 | 0.013444  | -0.304180 | 1.000000      | 0.003183  | 0.009612       | 0.014204        | -0.047820 |
| HasCrCard       | -0.005458   | -0.011721 | 0.022583  | -0.014858 | 0.003183      | 1.000000  | -0.011866      | -0.009933       | -0.007138 |
| IsActiveMember  | 0.025651    | 0.085472  | -0.028362 | -0.010084 | 0.009612      | -0.011866 | 1.000000       | -0.011421       | -0.156128 |
| EstimatedSalary | -0.001384   | -0.007201 | 0.007784  | 0.012797  | 0.014204      | -0.009933 | -0.011421      | 1.000000        | 0.012097  |
| Exited          | -0.027094   | 0.285323  | -0.014001 | 0.118533  | -0.047820     | -0.007138 | -0.156128      | 0.012097        | 1.000000  |

```
#Perform Bivariate Analysis
```

```
import statsmodels.api as sm
```

```

#define response variable y
= df['CreditScore']

#define explanatory variable x =
df[['EstimatedSalary']]

#add constant to predictor variables x
= sm.add_constant(x)

#fit linear regression model model
= sm.OLS(y, x).fit()

#view model summary print(model.summary())

```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          CreditScore   R-squared:                0.000
Model:                  OLS          Adj. R-squared:           -0.000
Method:                 Least Squares   F-statistic:              0.01916
Date:                  Thu, 29 Sep 2022   Prob (F-statistic):       0.890
Time:                  14:58:55         Log-Likelihood:          -59900.
No. Observations:      10000           AIC:                    1.198e+05
Df Residuals:          9998           BIC:                    1.198e+05
Df Model:               1
Covariance Type:       nonrobust
=====

```

|                 | coef       | std err  | t       | P> t  | [0.025    | 0.975]   |
|-----------------|------------|----------|---------|-------|-----------|----------|
| const           | 650.7617   | 1.940    | 335.407 | 0.000 | 646.958   | 654.565  |
| EstimatedSalary | -2.326e-06 | 1.68e-05 | -0.138  | 0.890 | -3.53e-05 | 3.06e-05 |

```

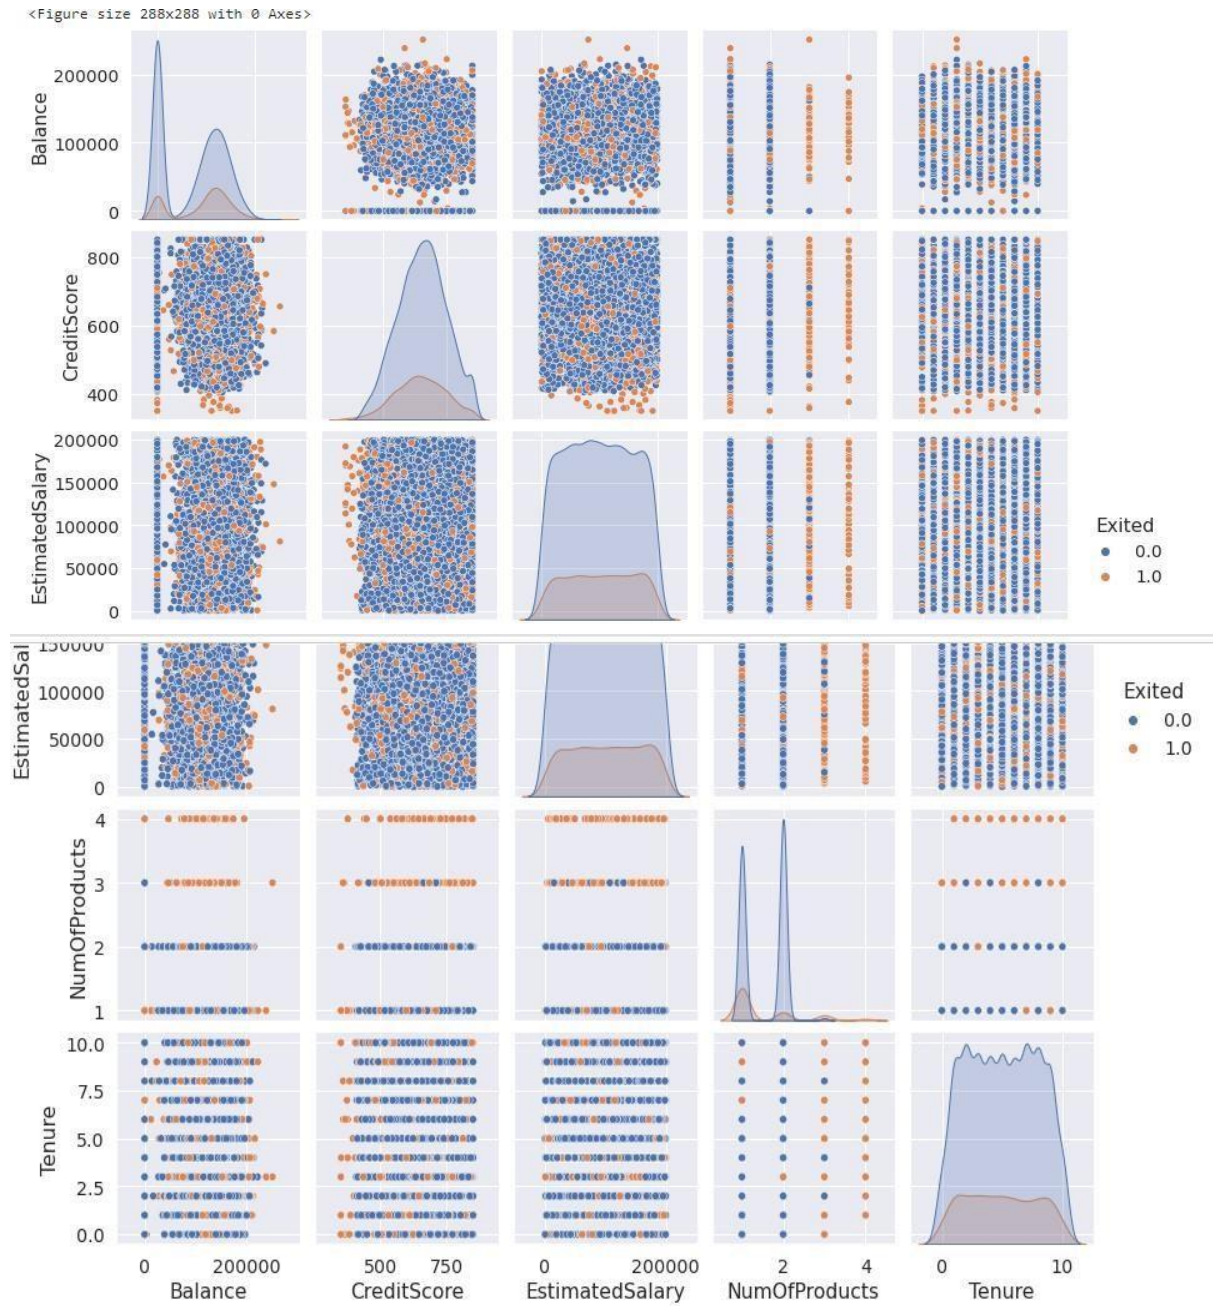
=====
Omnibus:                 132.939   Durbin-Watson:           2.014
Prob(Omnibus):            0.000   Jarque-Bera (JB):        84.242
Skew:                     -0.072   Prob(JB):                5.10e-19
Kurtosis:                 2.574   Cond. No.:               2.32e+05
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.32e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
#Perform Multivariate Analysis plt.figure(figsize=(4,4))
sns.pairplot(data=df[["Balance","CreditScore","EstimatedSalary","NumOfProducts",
"Tenure","Exited"]],hue="Exited")
```



### Question-3:

3. Perform descriptive statistics on the dataset.

#### Solution:

```
#Perform Descriptive Statistics
df=pd.DataFrame(df) print(df.sum())
```

```
CreditScore      6505288.0
Geography        FranceSpainFranceFranceSpainSpainFranceGermany...
Gender           FemaleFemaleFemaleFemaleFemaleMaleMaleFemaleMa...
Age              389218.0
Tenure           50128.0
Balance          764858892.88
NumOfProducts    15302.0
HasCrCard         7055.0
IsActiveMember    5151.0
EstimatedSalary  1000902398.81
Exited           2037.0
dtype: object
```

```
#Perform Descriptive Statistics print("----
Sum Value-----") print(df.sum(1)) print("-----
-----") print("----Product
Value-----") print(df.prod())
print("-----")
```

```
----Sum Value-----
0      102015.88
1      197002.44
2      274149.37
3       94567.63
4      205492.92
...
9995    97088.64
9996   159633.38
9997    42840.58
9998   168784.83
9999   169159.57
Length: 10000, dtype: float64
-----Product Value-----
CreditScore      inf
Age              inf
Tenure           0.0
Balance          0.0
NumOfProducts    inf
HasCrCard        0.0
IsActiveMember   0.0
EstimatedSalary  inf
Exited           0.0
dtype: float64
-----
```

```
#Perform Descriptive Statistics print("-----
-Mean Value-----") print(df.mean())
print("-----") print("-----
-----Median Value-----")
print(df.median()) print("-----
-----") print("-----Mode Value-----")
```



```

---") print(df.mode()) print("-----
-----")
-----Mean Value-----
CreditScore      650.528800
Age              38.921800
Tenure           5.012800
Balance          76485.889288
NumOfProducts    1.530200
HasCrCard        0.705500
IsActiveMember   0.515100
EstimatedSalary  100090.239881
Exited           0.203700
dtype: float64
-----
-----Median Value-----
CreditScore      652.000
Age              37.000
Tenure           5.000
Balance          97198.540
NumOfProducts    1.000
HasCrCard        1.000
IsActiveMember   1.000
EstimatedSalary  100193.915
Exited           0.000
dtype: float64
-----
-----Mode Value-----
   CreditScore Geography Gender   Age  Tenure  Balance  NumOfProducts
0         850.0    France   Male  37.0     2.0     0.0             1.0

   HasCrCard  IsActiveMember  EstimatedSalary  Exited
0         1.0             1.0         24924.92     0.0
-----

```

#### Question-4:

4.Handle the Missing values

#### Solution:

```

#Handling with missing Values df.isnull().values;
#Checking values are null

#Handling with missing Values df.notnull()#Checking
values are not null

```

Out[16]:

|      | CreditScore | Geography | Gender | Age  | Tenure | Balance | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSalary | Exited |
|------|-------------|-----------|--------|------|--------|---------|---------------|-----------|----------------|-----------------|--------|
| 0    | True        | True      | True   | True | True   | True    | True          | True      | True           | True            | True   |
| 1    | True        | True      | True   | True | True   | True    | True          | True      | True           | True            | True   |
| 2    | True        | True      | True   | True | True   | True    | True          | True      | True           | True            | True   |
| 3    | True        | True      | True   | True | True   | True    | True          | True      | True           | True            | True   |
| 4    | True        | True      | True   | True | True   | True    | True          | True      | True           | True            | True   |
| ...  | ...         | ...       | ...    | ...  | ...    | ...     | ...           | ...       | ...            | ...             | ...    |
| 9995 | True        | True      | True   | True | True   | True    | True          | True      | True           | True            | True   |
| 9996 | True        | True      | True   | True | True   | True    | True          | True      | True           | True            | True   |
| 9997 | True        | True      | True   | True | True   | True    | True          | True      | True           | True            | True   |
| 9998 | True        | True      | True   | True | True   | True    | True          | True      | True           | True            | True   |
| 9999 | True        | True      | True   | True | True   | True    | True          | True      | True           | True            | True   |


10000 rows × 11 columns

### Question-5:

5. Find the outliers and replace the outliers

#### Solution:

*#Find outliers & replace the outliers* `sns.boxplot(df['Balance'])`



```
#Find outliers & replace the outliers
print(np.where(df['Balance']>100000))

(array([ 2,  4,  5, ..., 9987, 9993, 9999]),)
```

In [19]:

```
#Find outliers & replace the outliers from
scipy import stats import numpy as np    z =
np.abs(stats.zscore(df["EstimatedSalary"]))
print(z)
```

```

0      0.021886
1      0.216534
2      0.240687
3      0.108918
4      0.365276
...
9995    0.066419
9996    0.027988
9997    1.008643
9998    0.125231
9999    1.076370
Name: EstimatedSalary, Length: 10000, dtype: float64

```

### Question-6:

6. Check for Categorical columns and perform encoding

#### Solution:

```

#Check for categorical columns & performs encoding from
sklearn.preprocessing import LabelEncoder

```

```

df['Gender'].unique() df['Gender'].value_counts()
encoding=LabelEncoder()
df["Gender"]=encoding.fit_transform(df.iloc[:,1].values) df

```

```

1: #Check for categorical columns & performs encoding
from sklearn.preprocessing import LabelEncoder
df['Gender'].unique()

```

```

1: array(['Female', 'Male'], dtype=object)

```

```

1: #Check for categorical columns & performs encoding
df['Gender'].value_counts()

```

```

1: Male      5457
   Female    4543
   Name: Gender, dtype: int64

```



```
Out[22]:
```

|      | CreditScore | Geography | Gender | Age  | Tenure | Balance   | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSalary | Exited |
|------|-------------|-----------|--------|------|--------|-----------|---------------|-----------|----------------|-----------------|--------|
| 0    | 619.0       | France    | 0      | 42.0 | 2.0    | 0.00      | 1.0           | 1.0       | 1.0            | 101348.88       | 1.0    |
| 1    | 608.0       | Spain     | 2      | 41.0 | 1.0    | 83807.86  | 1.0           | 0.0       | 1.0            | 112542.58       | 0.0    |
| 2    | 502.0       | France    | 0      | 42.0 | 8.0    | 159660.80 | 3.0           | 1.0       | 0.0            | 113931.57       | 1.0    |
| 3    | 699.0       | France    | 0      | 39.0 | 1.0    | 0.00      | 2.0           | 0.0       | 0.0            | 93826.63        | 0.0    |
| 4    | 850.0       | Spain     | 2      | 43.0 | 2.0    | 125510.82 | 1.0           | 1.0       | 1.0            | 79084.10        | 0.0    |
| ...  | ...         | ...       | ...    | ...  | ...    | ...       | ...           | ...       | ...            | ...             | ...    |
| 9995 | 771.0       | France    | 0      | 39.0 | 5.0    | 0.00      | 2.0           | 1.0       | 0.0            | 96270.64        | 0.0    |
| 9996 | 516.0       | France    | 0      | 35.0 | 10.0   | 57369.61  | 1.0           | 1.0       | 1.0            | 101699.77       | 0.0    |
| 9997 | 709.0       | France    | 0      | 36.0 | 7.0    | 0.00      | 1.0           | 0.0       | 1.0            | 42085.58        | 1.0    |
| 9998 | 772.0       | Germany   | 1      | 42.0 | 3.0    | 75075.31  | 2.0           | 1.0       | 0.0            | 92888.52        | 1.0    |
| 9999 | 792.0       | France    | 0      | 28.0 | 4.0    | 130142.79 | 1.0           | 1.0       | 0.0            | 38190.78        | 0.0    |

10000 rows x 11 columns

### Question-7:

7. Split the data into dependent and independent variables.

#### Solution:

```
#Split the data into Dependent & Independent Variables
print("-----Dependent Variables-----")
X=df.iloc[:,1:4] print(X) print("-----")
print("-----Independent Variables-----") Y=df.iloc[:,4] print(Y)
print("-----")
```

### Question-8:

8. Scale the independent variables

#### Solution:

```
#Split the data into Dependent & Independent Variables
print("-----Dependent Variables-----")
X=df.iloc[:,1:4] print(X) print("-----")
print("-----Independent Variables-----") Y=df.iloc[:,4]
print(Y)
print("-----")
```

### Question-9:

9. Split the data into training and testing

#### Solution:

```
#Split the data into training & testing from sklearn.model_selection
import train_test_split
```

In [34]:

```
#Split the data into training & testing
x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=4, random_state=4) x_train x_test y_train y_test
```

```
Out[31]:
```

|      | const | EstimatedSalary |
|------|-------|-----------------|
| 1603 | 1.0   | 23305.85        |
| 8713 | 1.0   | 41248.80        |
| 4561 | 1.0   | 143317.42       |
| 6600 | 1.0   | 174123.16       |

```
In [32]: #Split the data into training & testing
y_train
```

```
Out[32]:
```

|      |       |
|------|-------|
| 2558 | 727.0 |
| 7642 | 811.0 |
| 8912 | 623.0 |
| 3319 | 430.0 |
| 6852 | 600.0 |
| ...  |       |
| 456  | 733.0 |
| 6017 | 487.0 |
| 709  | 686.0 |
| 8366 | 637.0 |
| 1146 | 614.0 |

Name: CreditScore, Length: 9996, dtype: float64

Out[34]:

|      | const | EstimatedSalary |
|------|-------|-----------------|
| 2558 | 1.0   | 137903.54       |
| 7642 | 1.0   | 121765.00       |
| 8912 | 1.0   | 109470.34       |
| 3319 | 1.0   | 2923.61         |
| 6852 | 1.0   | 7312.25         |
| ...  | ...   | ...             |
| 456  | 1.0   | 7666.73         |
| 6017 | 1.0   | 9085.00         |
| 709  | 1.0   | 147794.63       |
| 8366 | 1.0   | 102515.42       |
| 1146 | 1.0   | 54776.64        |

9996 rows × 2 columns